

ARE 213 PS 1a

Aaron Watt

TOTAL POINTS

44 / 45

QUESTION 1

1 10 pts

1.1 a 5 / 5

✓ - 0 pts Correct

💬 Nice

1.2 b 5 / 5

✓ - 0 pts Correct

💬 Good

QUESTION 2

2 15 pts

2.1 a 5 / 5

✓ - 0 pts Correct

💬 Good

3.2 b 5 / 5

✓ - 0 pts Correct

💬 Good

3.3 C 4 / 5

- 0 pts Correct

- 1 Point adjustment

💬 Some of the other health conditions, eg heart condition, are shown to be caused by smoking and probably shouldn't be included. Things that affect the outcome but are uncorrelated with treatment should be excluded because they effectively introduce noise.

3.4 d 5 / 5

✓ - 0 pts Correct

💬 Good

2.2 b 5 / 5

✓ - 0 pts Correct

💬 Good

2.3 C 5 / 5

✓ - 0 pts Correct

💬 Good

QUESTION 3

3 20 pts

3.1 a 5 / 5

✓ - 0 pts Correct

💬 Good

1.*Let it Snow!

Before getting started with the data work, first consider the table from Snow (1855) reproduced in the lecture notes (“Snow’s Table IX”). The table reports only means.

- (a) Develop an approximate 95% confidence interval for “Deaths per 10,000 Houses” for Southwark and Vauxhall customers. Develop another 95% CI for the same quantity for Lambeth. Do the confidence intervals overlap?

TABLE IX.

	Number of houses.	Deaths from Cholera.	Deaths in each 10,000 houses.
Southwark and Vauxhall Company	40,046	1,263	315
Lambeth Company	26,107	98	37
Rest of London	256,423	1,422	59

Figure 1: The orginal table from *On the Mode of Communication of Cholera* (John Snow, 1855). Note the rate of deaths per 10,000 houses for the Rest of London is incorrect based on the first two columns.
 $1,422/256,423 \cdot 10,000 \approx 55.5$ not 59.

Assuming that, within a region $k \in \{a, b, c\}$, the indicator of one person in household i dying from Cholera X_i^k is distributed as a Bernoulli random variable with probability $p_k - X_i^k \stackrel{i.i.d.}{\sim} \text{Bernoulli}(p_k)$. Then deaths per 10,000 households in region k , R_k , is a function of the sum of i.i.d. Bernoulli random variables. Let N_k be the number of houses in region k , then

$$R_k = \frac{10,000}{N_k} \sum_{i=1}^{N_k} X_i^k \sim \frac{10,000}{N_k} \text{Binom}(N_k, p_k)$$

Let $Z^k \sim \text{Binom}(N_k, p_k)$. Then the variance of R_k is

$$\begin{aligned} \text{var}(R_k) &= \text{var}\left(\frac{10,000}{N_k} Z^k\right) \\ &= \left(\frac{10,000}{N_k}\right)^2 \text{var}(Z^k) \\ &= \left(\frac{10,000}{N_k}\right)^2 N_k \cdot p_k \cdot (1 - p_k) \\ &= \frac{(10,000)^2}{N_k} p_k (1 - p_k) \end{aligned}$$

Let z_α be the z-score at significance level α . Then, for large enough N_k (which 26,000 is probably large enough), the 95% ($\alpha = (1 - 95\%)/2$) confidence interval for R_k would be:

$$\begin{aligned}\text{CI}_{95} &= R_k \pm z_\alpha \sqrt{\text{var}(R_k)} \\ &= R_k \pm z_\alpha \sqrt{\frac{(10,000)^2}{N_k} p_k(1 - p_k)} \\ &= R_k \pm 10,000 z_\alpha \sqrt{\frac{p_k(1 - p_k)}{N_k}}\end{aligned}$$

And an unbiased estimator for the probability of death p_k would be the mean of the indicators:

$$\hat{p}_k = \frac{1}{N_k} \sum_{i=1}^{N_k} X_i^k = \frac{D_k}{N_k} = \bar{X}_i^k$$

where D_k is the number of deaths in region k (the sum of the X_i^k 's).

Calculating \hat{p}_k from the Snow table, we get

k	R_k	CI_k lower	CI_k upper
a	315.39	298.27	332.50
b	37.54	30.12	44.96
c	55.46	52.58	58.33

No, the confidence intervals for regions a and b do not overlap.

1.1 a 5 / 5

✓ - 0 pts Correct

 Nice

- (b) Discuss either formally or intuitively the critical assumption that underlies your confidence intervals. Give a 2 or 3 sentence quote from Snow's description (re-produced in Freedman (1991)) that supports this assumption.

The critical assumption of the above calculation is the the observations in each of the regions k are statistically independent of each other – that the probability of death in one region is not affected by the probability or realization of death in the other region, or by the underlying characteristics of the population in the different regions. Because the different regions were being supplied from different water sources, then conditional on the upstream water supply not being contaminated, if we assume the water was the source of Cholera infection, then the two groups death rates should be independent.

To support this idea, this section, from pg 75, of Snow's 1955 book is of interest:

Each Company supplies both rich and poor, both large houses and small; there is no difference either in the condition or occupation of the persons receiving the water of the different Companies...

... The experiment, too, was on the grandest scale. No fewer than three hundred thousand people of both sexes, of every age and occupation, and of every rank and station, from gentlefolks down to the very poor, were divided into two groups without their choice, and, in most cases, without their knowledge; one group being supplied with water containing the sewage of London, and, amongst it, whatever might have come from the cholera patients, the other group having water quite free from such impurity.

This natural experiment provided evidence that, even when the groups are essentially identical in distributions of characteristics and treatment of water supply is as good as random, that the group getting the water from downstream was statistically significantly more likely to die from Cholera.

We now move to some analysis of real data. The data portions of Problem Sets 1a and 1b are based heavily on the paper Almond, Chay, and Lee (2005), and problem sets from Ken Chay and John DiNardo based on some of the data used in the paper. The goal of this assignment is to examine the research question: what is the causal effect of maternal smoking during pregnancy on infant birthweight and other infant health outcomes. The data for the problem set is an extract of all births from the 1993 National Natality Detail Files for Pennsylvania. Each observation represents an infant-mother match. The data in Stata format can be downloaded from the bCourses website. There should be 48 variables in the data and, after you are finished with the cleaning steps described below, 114,610 observations.

The data here are “real” and quite imperfect, which will help simulate the unpleasantness of real world data work. Unlike the real world where you will confront this bleak situation largely alone, I will provide you with some hints for working your way through the raw data. You can download part of the codebook for the data to help you figure out the relevant variables.

1.2 b 5 / 5

✓ - 0 pts Correct

 Good

2. Real Data

The first order of business is to go through the code book, decide on the relevant variables, and process the data. This involves several steps:

- (a) Fix missing values. In the the data set several variables take on a value of, say, 9999 if missing. We have already checked for missing observations for about 2/3 of the variables. The remaining variables need to be checked and are the last 15 in the variable list (i.e. from 'cardiac' to 'wgain'). Refer to the codebook for missing value codes. Produce an analysis data set that drops any observation with missing values.

Table 1: Missing values, by variable name.

varname	missing_code	num_missing
herpes	8	6
tobacco	9	65
cigar	99	1077
cigar6	6	1077
alcohol	9	105
drink	99	2116
drink5	5	2116
wgain	99	3048

The observations with missing values have been removed, bringing our number of observations down from 120,461 to 114,610 (5851 observations dropped).

2.1 a 5 / 5

✓ - 0 pts Correct

 Good

- (b) If this were a real research project you would want to consider other approaches to missing data besides termination with extreme prejudice. What observations do you have to drop because of missing data. Might this affect your results? Do the data appear to be missing completely at random? How might you assess whether the data appear to be missing at random?

The data do not appear to be missing completely at random. MCAR (Missing Completely at Random) implies that there is no relationship between the missing observations and the observable variables, but we can see from figure 2 that the distribution of cigarette smoking is different between observations that have missing drink and those that are not missing drink. We also see that most observations missing tobacco are also missing drink.

Another crude method to see this is to create a single dummy variable for missingness in any variable. Then, using the dummy to create two sub-samples, one which has complete data, and another with any missingness, we can conduct a series of t-tests (with 5% significance level) on whether means of each variables are different across the two sub-samples. We reject null of equivalence on 37 out of 48 variables, giving another evidence that missing is not completely random.

We could further explore missing relationships by running regressions of the variables on the missingness indicators for several other variables, then testing if the coefficient on the indicators is significantly different from zero. We would need to be a bit more precise, but this could produce a test of significance of a relationship with the other variables. This would at least test for MCAR.

The problem with testing MAR is that the assumption is violated when missingness depends on unobservables. This is inherently hard to test.

Missing data matrix

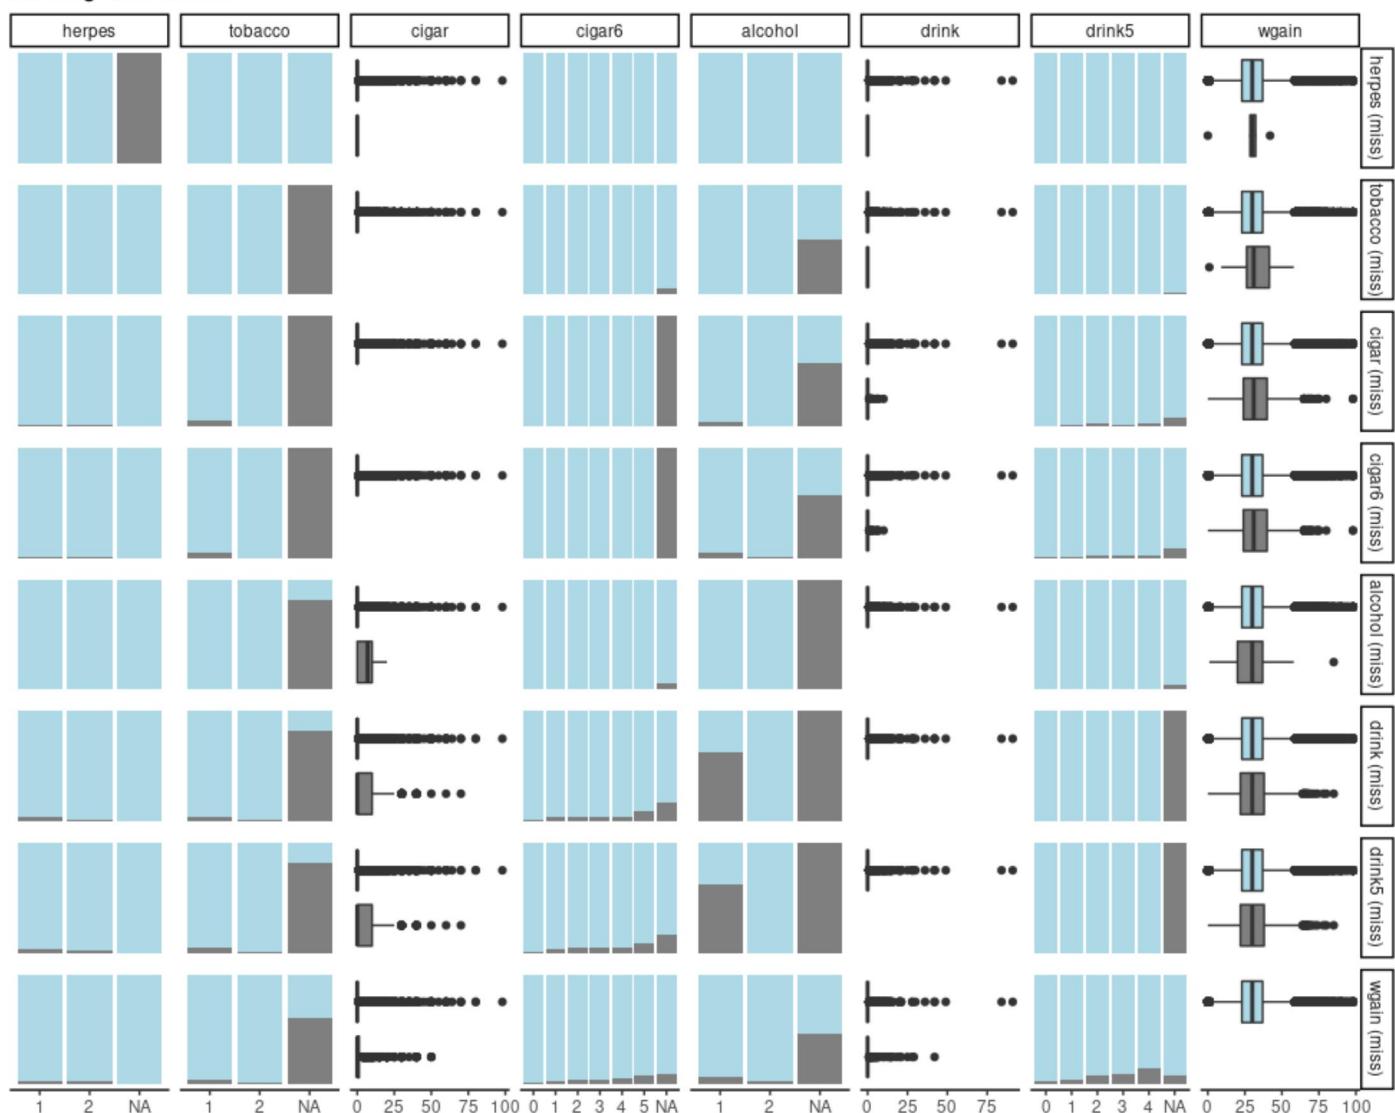


Figure 2: A plot to show the distribution of variables, broken out by missing and not missing in other variables.

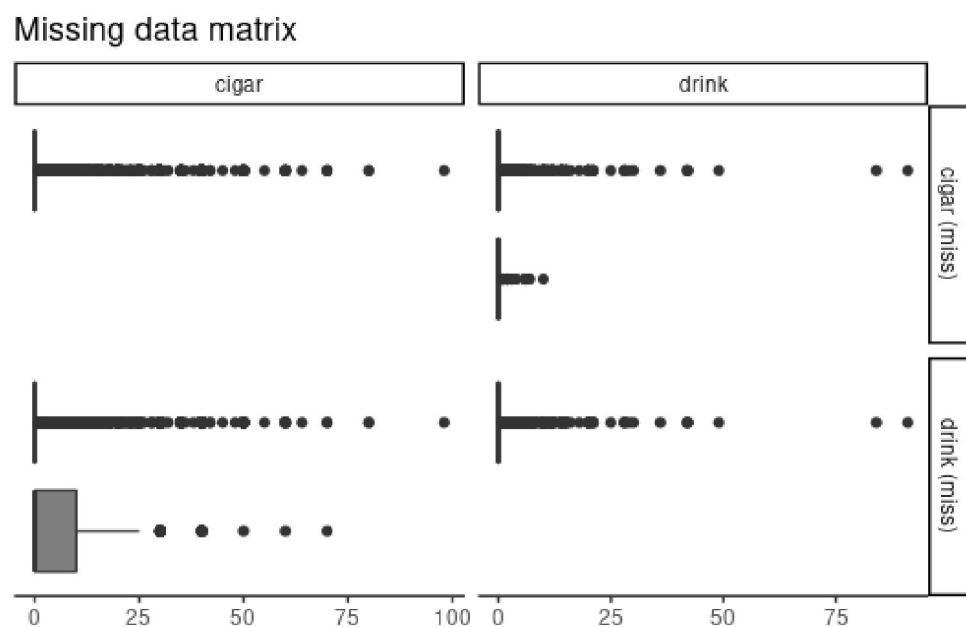


Figure 3: Looking closer at the missingness relationship between drinking and smoking.

2.2 b 5 / 5

✓ - 0 pts Correct

 Good

(c) Produce a summary table describing the final analysis data set.

Table 2: Table 1: Summary Statistics

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Pctl(75)	Max
rectype	114,610	1.262	0.440	1	1	2	2
pldel3	114,610	1.018	0.133	1	1	1	2
birattnnd	114,610	1.202	0.564	1	1	1	5
cntocpop	114,610	1.443	1.137	0	0	2	3
stresfip	114,610	41.743	2.167	0	42	42	55
dimage	114,610	27.757	5.699	12	24	32	49
ormoth	114,610	0.091	0.522	0	0	0	5
mrace3	114,610	1.259	0.657	1	1	1	3
dmeduc	114,610	13.211	2.272	0	12	16	17
dmar	114,610	1.251	0.434	1	1	2	2
adequacy	114,610	1.297	0.546	1	1	2	3
nlnbnl	114,610	0.967	1.148	0	0	1	12
dlivord	114,610	1.986	1.174	1	1	2	14
dtotord	114,610	2.420	1.520	1	1	3	24
totord9	114,610	2.407	1.458	1	1	3	8
monpre	114,610	2.502	1.326	0	2	3	9
nprevist	114,610	11.153	3.524	0	9	13	49
disllb	114,610	350.412	362.325	0	31	777	777
isllb10	114,610	3.321	3.188	0	0	6	9
dfage	114,610	30.062	6.410	13	26	34	78
orfath	114,610	0.095	0.531	0	0	0	5
dfeduc	114,610	13.277	2.325	0	12	16	17
birmon	114,610	6.474	3.394	1	4	9	12
weekday	114,610	4.047	1.881	1	2	6	7
dgestat	114,610	39.153	2.445	17	38	40	47
csex	114,610	1.485	0.500	1	1	2	2
dbrwt	114,610	3,373.291	585.175	227	3,062	3,742	6,067
dplural	114,610	1.028	0.174	1	1	1	4
omaps	114,610	8.117	1.260	0	8	9	10
fmaps	114,610	9.009	0.707	0	9	9	10
clingest	114,610	39.109	2.057	17	38	40	44
delmeth5	114,610	1.549	1.010	1	1	1	5
anemia	114,610	1.990	0.099	1	2	2	2
cardiac	114,610	1.993	0.083	1	2	2	2
lung	114,610	1.993	0.085	1	2	2	2
diabetes	114,610	1.973	0.162	1	2	2	2
herpes	114,610	1.994	0.078	1	2	2	2
chyper	114,610	1.992	0.087	1	2	2	2
phyper	114,610	1.969	0.172	1	2	2	2
pre4000	114,610	1.986	0.119	1	2	2	2

preterm	114,610	1.986	0.118	1	2	2	2
tobacco	114,610	1.841	0.366	1	2	2	2
cigar	114,610	1.907	5.297	0	0	0	98
cigar6	114,610	0.346	0.861	0	0	0	5
alcohol	114,610	1.990	0.098	1	2	2	2
drink	114,610	0.031	0.619	0	0	0	91
drink5	114,610	0.020	0.230	0	0	0	4
wgain	114,610	30.356	11.884	0	23	37	98

2.3 C 5 / 5

✓ - 0 pts Correct

 Good

3. Smoking and birth weight

The next part of the assignment is to try to estimate the “causal” effect of maternal smoking during pregnancy on infant birth weight. Let’s start out using techniques that are familiar, and think about whether they are likely to work in this context. Answer the following questions.

- (a) Compute the mean difference in APGAR scores (both five and one minute versions) as well as birthweight by smoking status.

Denote the group with tobacco usage during pregnancy as group 1, and without tobacco usage as group 2. On average, group 2 has 0.043 and 0.013 higher APGAR scores on one and five minute versions respectively. A simple t-test demonstrates they are significant at 1 and 5% significance levels, respectively. Furthermore, group 2 has average birth-weight 246 grams higher, significant at 1% significance level.

	Mean_Smoker	Mean_NonSmoker	Diff_Means	statistic	p.value	95% conf.int
1m APGAR	8.072	8.115	-0.043	-4.300	0.000	[-0.062, -0.024]
5m APGAR	8.992	9.005	-0.013	-2.311	0.023	[-0.024, -0.002]
BWHT	3160.306	3407.223	-246.917	-55.214	0.000	[-255.683, -238.152]

3.1 a 5 / 5

✓ - 0 pts Correct

 Good

- (b) Under what circumstances can one identify the average treatment effect of maternal smoking by comparing the unadjusted difference in mean birth weight of infants of smoking and non-smoking mothers? Estimate its impact under this assumption. Provide and comment on some evidence for or against the validity of the assumption (A useful “Table 1” of any paper is one that describes the overall averages of the observations, and then describes the subsets of people who do and do not receive the treatment (when it is binary)).

If assignment of maternal smoking was completely random, the differences in means that was reported in Q3-Part(a) would be the ATE of maternal smoking on infant birth-weight. Then, just as in any RCTs, we would see balance across two-groups on any observable variables measured before the ‘hypothetical assignment’ of maternal smoking. Assume treatment of smoking during pregnancy were assigned in the very early stages of pregnancy, then we must observe balance across the ‘treatment’ and ‘control’ groups of covariates either fixed (e.g. age, race) or measured before pregnancy (e.g. chronic conditions), which we deemed amounts to 15 variables in the data-set. However, the balance table below suggests that balance is generally not reached across 9 of 15 covariates based on indicator variable for maternal smoking, which clearly indicates non-random smoking patterns.

tobacco Variable		N	1 Mean	SD	N	2 Mean	SD	Test
dimage		18266	26.173	5.606	96344	28.057	5.667	F=1702.046***
ormoth		18266	0.064	0.449	96344	0.096	0.535	F=58.122**
mrace3		18266	1.258	0.667	96344	1.259	0.655	F=0.077
dmar		18266	1.482	0.5	96344	1.207	0.405	F=6516.468***
nlblnl		18266	1.152	1.224	96344	0.932	1.13	F=569.708***
dliword		18266	2.181	1.265	96344	1.949	1.153	F=599.644***
dfage		18266	28.961	6.646	96344	30.271	6.343	F=644.891***
orfath		18266	0.082	0.497	96344	0.097	0.537	F=13.405***
dfeduc		18266	12.128	1.672	96344	13.494	2.368	F=5555.681***
csex		18266	1.482	0.5	96344	1.486	0.5	F=1.072
cardiac		18266	1.994	0.078	96344	1.993	0.084	F=2.03
lung		18266	1.99	0.098	96344	1.993	0.082	F=18.421***
diabetes		18266	1.973	0.162	96344	1.973	0.162	F=0
herpes		18266	1.993	0.081	96344	1.994	0.077	F=1.06
chyper		18266	1.993	0.081	96344	1.992	0.089	F=3.713*

Statistical significance markers: * p<0.1; ** p<0.05; *** p<0.01

3.2 b 5 / 5

✓ - 0 pts Correct

 Good

- (c) Suppose that maternal smoking is randomly assigned conditional on the other observable “predetermined” determinants of infant birth weight. First discuss which (if any) of the variables contained in the data set can clearly be considered to be predetermined. In general, what kinds of variables can be considered predetermined and what kinds of variables cannot?

Generally, predetermined variables are those which are determined prior to the treatment period—in this case, prior to the decision to smoke. These are most obviously ascriptive characteristics such as race and sex, but could also variables that are clearly uncorrelated with treatment. In our case, these would be race (mrace3), sex (csex), heart condition (cardiac), diabetes (diabetes), herpes (herpes), and hypertension (chyper).

The variable that most obviously cannot be considered ”predetermined” is lung issues (lung) since smoking has been proven to directly cause lung issues. Other variables like education (dfeduc) are somewhere in between predetermined and not.

When running a regression, it is particularly important to control for the variables that are correlated with both treatment and birth weight if we want to isolate the causal impact of smoking on birth weight.

3.3 C 4 / 5

- **0 pts** Correct

- **1 Point adjustment**

- Some of the other health conditions, eg heart condition, are shown to be caused by smoking and probably shouldn't be included. Things that affect the outcome but are uncorrelated with treatment should be excluded because they effectively introduce noise.

- (d) What does “selection on observables” imply about the relationship between maternal smoking and unobservable determinants of birth weight conditional on the observables? Use a basic linear regression model, in conjunction with your answer to part (c), to estimate the impact of smoking and report your estimates. Under what circumstances is the average treatment effect identified?

Selection on observables implies that maternal smoking is conditionally independent of unobserved factors (conditional on the X 's we have observed).

If the covariates we included in our regression include all possible (relevant) covariates that are correlated with both smoking and birth weight and if the covariates of the mothers across all smoking categories share the same support, then the relationship from a regression of birth weight on smoking status will give us the true ATE of smoking on baby weight.

Below regressions suggest that smoking decreases birthweight by approximately 218 grams with statistical significance at 1% significance; smoking increases 1 minute and 5 minute APGAR scores, but this is statistically insignificant at all usual significance levels.

Table 3: Linear Regression Selection on Observables

	<i>Dependent variable:</i>		
	dbrwt	omaps	fmaps
	(1)	(2)	(3)
tobacco	217.986*** (4.753)	-0.011 (0.011)	-0.010 (0.006)
rectype	-47.944*** (3.867)	-0.044*** (0.009)	-0.016*** (0.005)
pldel3	142.614*** (14.231)	0.182*** (0.032)	0.208*** (0.018)
dmage	-4.559*** (0.527)	-0.007*** (0.001)	-0.003*** (0.001)
ormoth	-31.493*** (4.153)	-0.032*** (0.009)	-0.019*** (0.005)
mrace3	-105.103*** (2.929)	-0.074*** (0.007)	-0.046*** (0.004)
delmeth5	-10.360*** (1.643)	-0.098*** (0.004)	-0.032*** (0.002)
diabetes	-65.834*** (10.190)	0.165*** (0.023)	0.067*** (0.013)
chyper	264.745*** (18.750)	0.266*** (0.042)	0.144*** (0.024)
21 other covariates not shown			
Constant	1,558.841*** (88.273)	6.942*** (0.199)	8.099*** (0.112)
Observations	114,610	114,610	114,610
R ²	0.109	0.021	0.024
Adjusted R ²	0.109	0.021	0.024
Residual Std. Error (df = 114579)	552.347	1.246	0.699
F Statistic (df = 30; 114579)	468.596***	81.881***	93.051***

*p<0.1; **p<0.05; ***p<0.01

3.4 d 5 / 5

✓ - 0 pts Correct

 Good